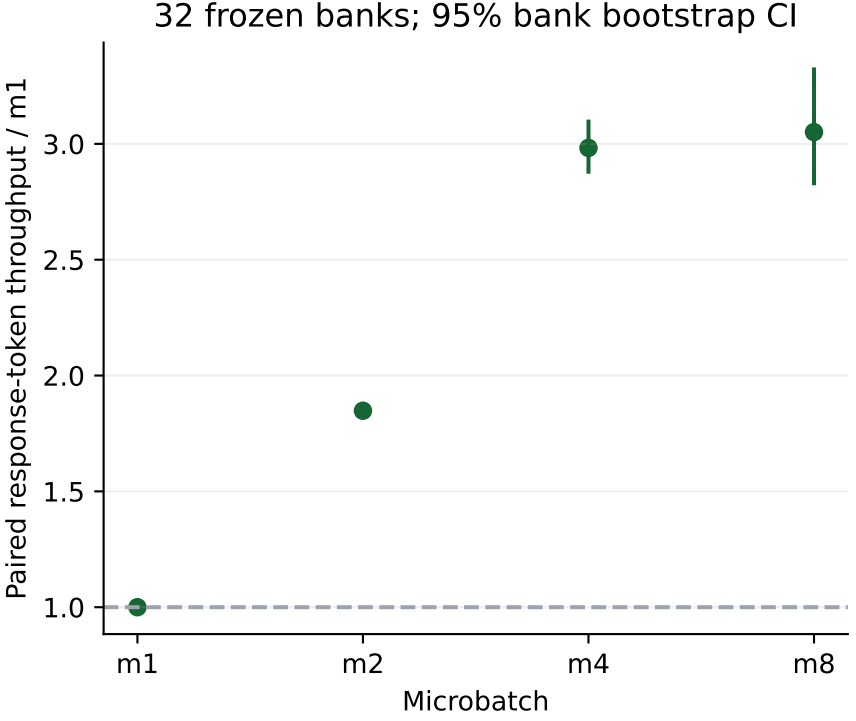
32 frozen banks; 95% bank bootstrap CI
Paired response-token throughput / m1
1.0
1.5
2.0
2.5
3.0
m1
m2
m4
m8
Microbatch

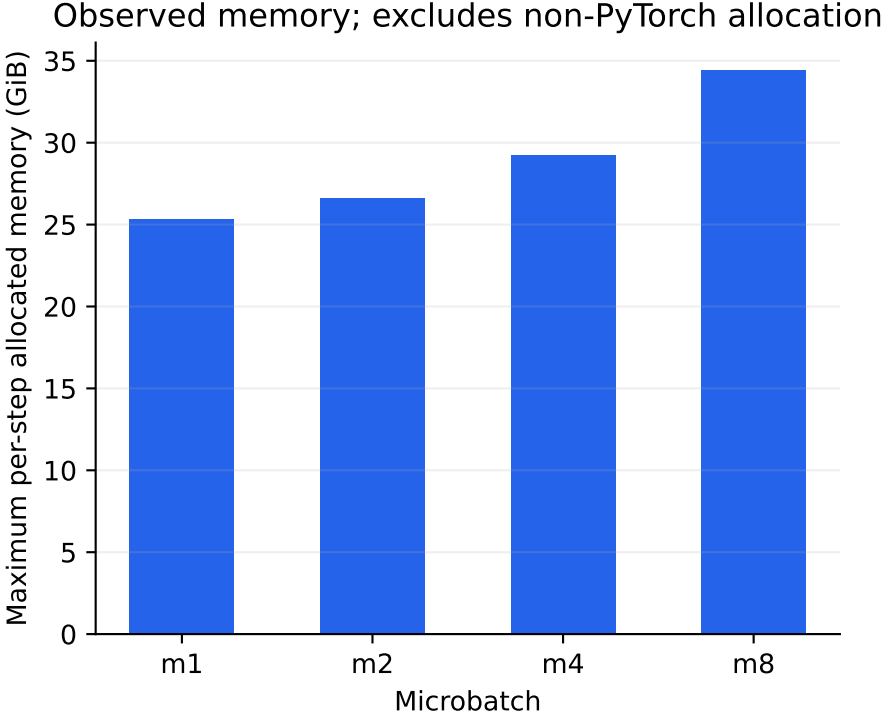
Observed memory; excludes non-PyTorch allocation
Maximum per-step allocated memory (GiB)
0
5
10
15
20
25
30
35
m1
m2
m4
m8
Microbatch